

Section 2: Integer & Floating Point Numbers

- Representation of integers: unsigned and signed
 - Unsigned and signed integers in C
 - Arithmetic and shifting
 - Sign extension
-
- Background: fractional binary numbers
 - IEEE floating-point standard
 - Floating-point operations and rounding
 - Floating-point in C

Unsigned Integers

- **Unsigned values are just what you expect**

- $b_7b_6b_5b_4b_3b_2b_1b_0 = b_72^7 + b_62^6 + b_52^5 + \dots + b_12^1 + b_02^0$
- Useful formula: $1+2+4+8+\dots+2^{N-1} = 2^N - 1$

- **You add/subtract them using the normal “carry/borrow” rules, just in binary**

$$\begin{array}{r}
 63 \qquad 00111111 \\
 + \underline{8} \qquad + \underline{00001000} \\
 71 \qquad 01000111
 \end{array}$$

Signed Integers

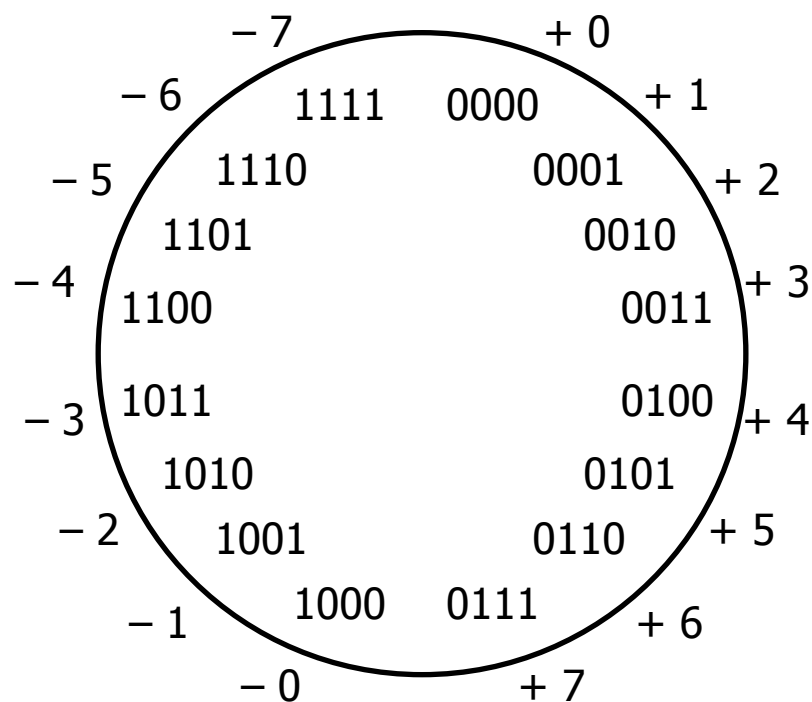
- **Let's do the natural thing for the positives**
 - They correspond to the unsigned integers of the same value
 - Example (8 bits): $0x00 = 0$, $0x01 = 1$, ..., $0x7F = 127$
- **But, we need to let about half of them be negative**
 - Use the high order bit to indicate *negative*: call it the “sign bit”
 - Call this a “sign-and-magnitude” representation
 - Examples (8 bits):
 - $0x00 = 00000000_2$ is non-negative, because the sign bit is 0
 - $0x7F = 01111111_2$ is non-negative
 - $0x85 = 10000101_2$ is negative
 - $0x80 = 10000000_2$ is negative...

Sign-and-Magnitude Negatives

■ How should we represent -1 in binary?

- Sign-and-magnitude: 10000001_2

Use the MSB for + or -, and the other bits to give magnitude



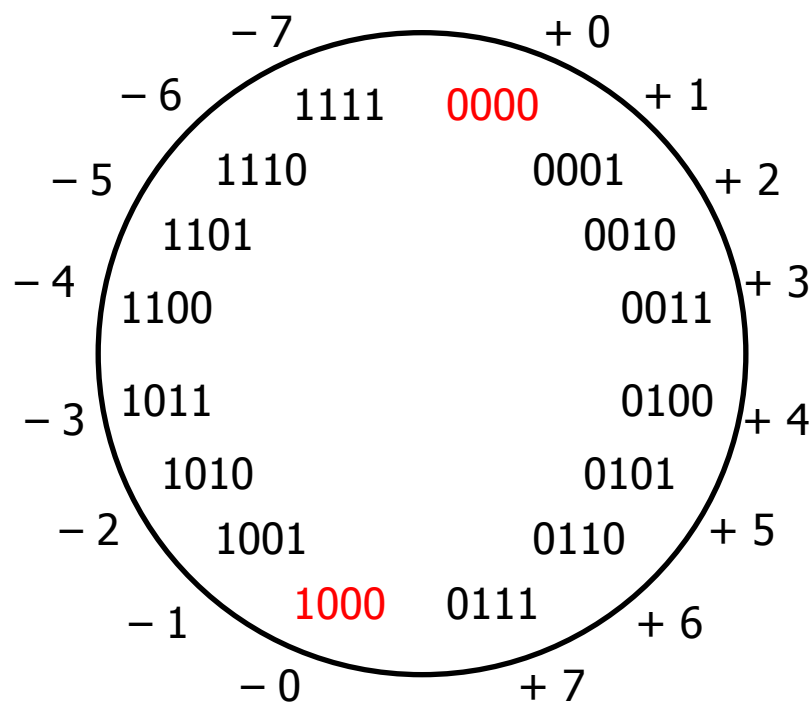
Sign-and-Magnitude Negatives

■ How should we represent -1 in binary?

- Sign-and-magnitude: 10000001_2

Use the MSB for + or -, and the other bits to give magnitude

(Unfortunate side effect: there are two representations of 0!)



Sign-and-Magnitude Negatives

■ How should we represent -1 in binary?


- Sign-and-magnitude: 10000001_2

Use the MSB for + or -, and the other bits to give magnitude
(Unfortunate side effect: there are two representations of 0!)

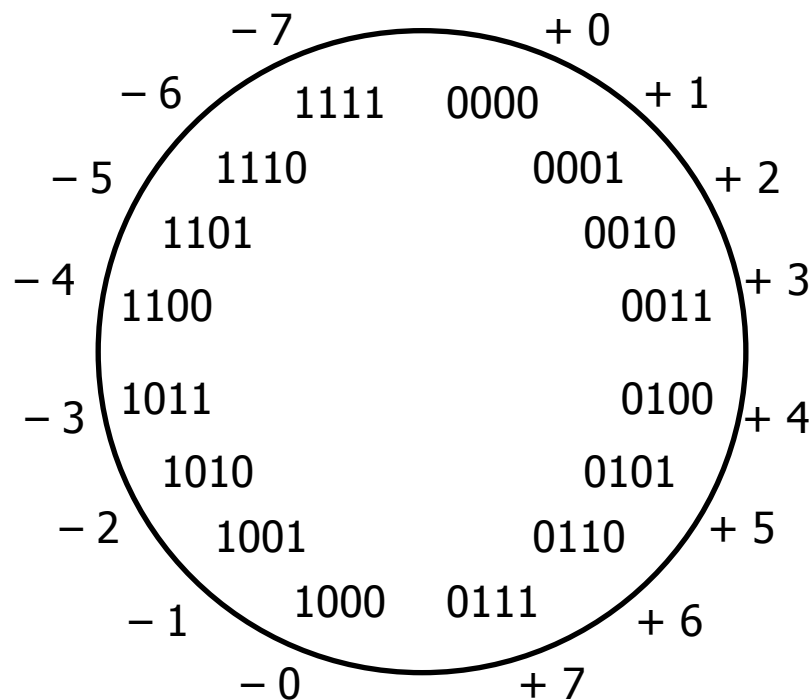
- Another problem: math is cumbersome

- Example:

$$4 - 3 \neq 4 + (-3)$$



0100
+ 1011
1111



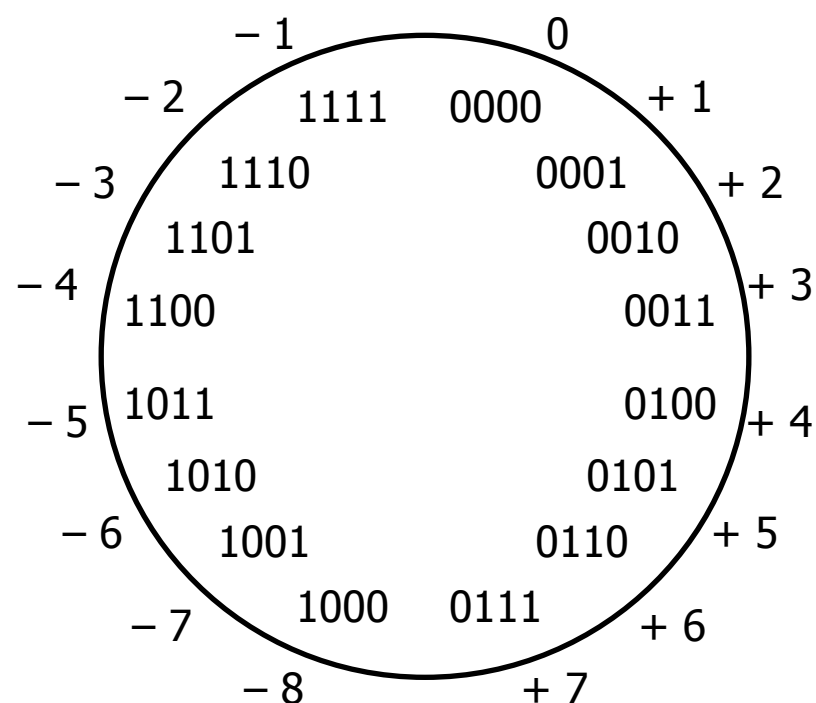
Two's Complement Negatives

■ How should we represent -1 in binary?

- Rather than a sign bit, let MSB have same value, but *negative* weight
 - W-bit word: Bits 0, 1, ..., W-2 add $2^0, 2^1, \dots, 2^{W-2}$ to value of integer when set, but bit W-1 adds -2^{W-1} when set
 - e.g. unsigned 1010_2 : $1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = 10_{10}$
 2's comp. 1010_2 : $-1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = -6_{10}$

- So -1 represented as 1111_2 ; all negative integers still have MSB = 1
- Advantages of two's complement: only one zero, simple arithmetic
- To get negative representation of any integer, take bitwise complement and then add one!

$$\sim x + 1 = -x$$



Two's Complement Arithmetic

- The same addition procedure works for both unsigned and two's complement integers
 - Simplifies hardware: only one adder needed
 - Algorithm: simple addition, discard the highest carry bit
 - Called “modular” addition: result is sum *modulo* 2^w
- Examples:

4 0100	4 0100	− 4 1100
+ 3 + 0011	− 3 + 1101	+ 3 + 0011
= 7 = 0111	= 1 1 0001	− 1 1111
	drop carry = 0001	

Two's Complement

■ Why does it work?

- Put another way: given the bit representation of a positive integer, we want the negative bit representation to always sum to 0 (ignoring the carry-out bit) when added to the positive representation
- This turns out to be the *bitwise complement plus one*
 - What should the 8-bit representation of -1 be?

$$\begin{array}{r}
 00000001 \\
 + \underline{????????} \\
 \hline
 00000000
 \end{array}
 \quad \text{(we want whichever bit string gives the right result)}$$

$$\begin{array}{r}
 00000010 \\
 + \underline{????????} \\
 \hline
 00000000
 \end{array}
 \qquad
 \begin{array}{r}
 00000011 \\
 + \underline{????????} \\
 \hline
 00000000
 \end{array}$$

Two's Complement

■ Why does it work?

- Put another way: given the bit representation of a positive integer, we want the negative bit representation to always sum to 0 (ignoring the carry-out bit) when added to the positive representation
- This turns out to be the *bitwise complement plus one*
 - What should the 8-bit representation of -1 be?

$$\begin{array}{r}
 00000001 \\
 + 11111111 \\
 \hline
 100000000
 \end{array}
 \quad \text{(we want whichever bit string gives the right result)}$$

$$\begin{array}{r}
 00000010 \\
 + \text{????????} \\
 \hline
 00000000
 \end{array}
 \qquad
 \begin{array}{r}
 00000011 \\
 + \text{????????} \\
 \hline
 00000000
 \end{array}$$

Two's Complement

■ Why does it work?

- Put another way: given the bit representation of a positive integer, we want the negative bit representation to always sum to 0 (ignoring the carry-out bit) when added to the positive representation

- This turns out to be the *bitwise complement plus one*

- What should the 8-bit representation of -1 be?

00000001

+11111111

100000000

(we want whichever bit string gives the right result)

00000010

+11111110

100000000

00000011

+11111101

100000000

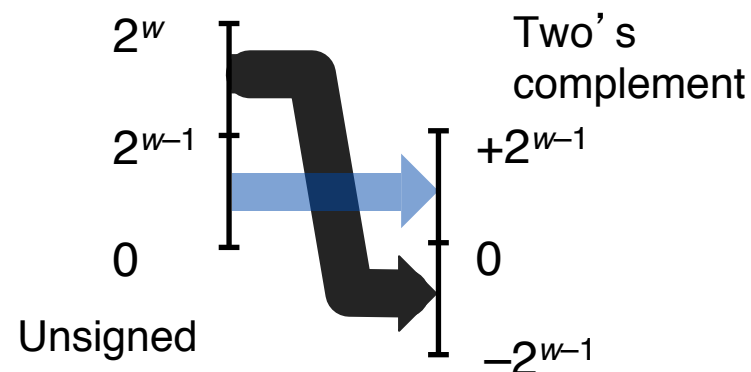
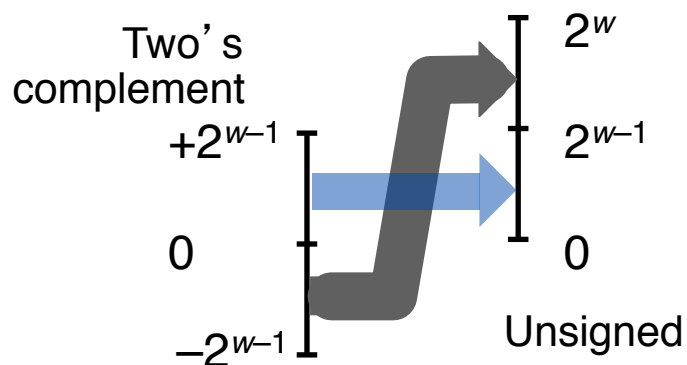
Unsigned & Signed Numeric Values

X	Unsigned	Signed
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

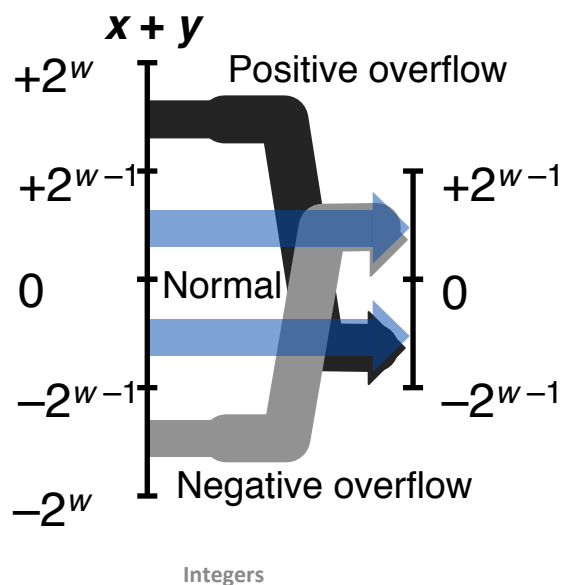
- Both signed and unsigned integers have limits
 - If you compute a number that is too big, you wrap: $6 + 4 = ?$ $15U + 2U = ?$
 - If you compute a number that is too small, you wrap: $-7 - 3 = ?$ $0U - 2U = ?$
- The CPU may be capable of “throwing an exception” for overflow on signed values
 - But it won't for unsigned
- C and Java just cruise along silently when overflow occurs...

Visualizations

- Same W bits interpreted as signed vs. unsigned:



- Two's complement (signed) addition: x and y are W bits wide



Values To Remember

■ Unsigned Values

- UMin = 0
 - 000...0
- UMax = $2^w - 1$
 - 111...1

■ Two's Complement Values

- TMin = -2^{w-1}
 - 100...0
- TMax = $2^{w-1} - 1$
 - 011...1
- Negative 1
 - 111...1 0xFFFFFFFF (32 bits)

Values for $W = 16$

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000