

Section 2: Integer & Floating Point Numbers

- Representation of integers: unsigned and signed
 - Unsigned and signed integers in C
 - Arithmetic and shifting
 - Sign extension
-
- Background: fractional binary numbers
 - IEEE floating-point standard
 - Floating-point operations and rounding
 - Floating-point in C

IEEE Floating Point

■ Analogous to scientific notation

- Not 12000000 but 1.2×10^7 ; not 0.0000012 but 1.2×10^{-6}
 - (write in C code as: 1.2e7; 1.2e-6)

■ IEEE Standard 754

- Established in 1985 as uniform standard for floating point arithmetic
 - Before that, many idiosyncratic formats
- Supported by all major CPUs today

■ Driven by numerical concerns

- Standards for handling rounding, overflow, underflow
- Hard to make fast in hardware but numerically well-behaved

Floating Point Representation

■ Numerical form:

$$V_{10} = (-1)^s * M * 2^E$$

- Sign bit **s** determines whether number is negative or positive
- Significand (mantissa) **M** normally a fractional value in range [1.0,2.0)
- Exponent **E** weights value by a (possibly negative) power of two

■ Representation in memory:

- MSB **s** is sign bit **s**
- **exp** field encodes **E** (but is *not equal* to E)
- **frac** field encodes **M** (but is *not equal* to M)




Precisions

- Single precision: 32 bits




- Double precision: 64 bits



Normalization and Special Values


$$V = (-1)^S * M * 2^E$$



- **“Normalized” means the mantissa M has the form 1.xxxxx**
 - 0.011×2^5 and 1.1×2^3 represent the same number, but the latter makes better use of the available bits
 - Since we know the mantissa starts with a 1, we don't bother to store it
- **How do we represent 0.0? Or special / undefined values like 1.0/0.0?**

Normalization and Special Values


$$V = (-1)^S * M * 2^E$$



- **“Normalized” means the mantissa M has the form 1.xxxxx**
 - 0.011×2^5 and 1.1×2^3 represent the same number, but the latter makes better use of the available bits
 - Since we know the mantissa starts with a 1, we don't bother to store it
- **Special values:**
 - The bit pattern 00...0 represents **zero**
 - If **exp** == 11...1 and **frac** == 00...0, it represents ∞
 - e.g. $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -1.0/0.0 = -\infty$
 - If **exp** == 11...1 and **frac** != 00...0, it represents **NaN**: “Not a Number”
 - Results from operations with undefined result, e.g. $\text{sqrt}(-1)$, $\infty - \infty$, $\infty * 0$

Normalized Values


$$V = (-1)^S * M * 2^E$$



- **Condition:** $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$
- **Exponent coded as *biased* value:** $E = \text{exp} - \text{Bias}$
 - exp is an *unsigned* value ranging from 1 to $2^k - 2$ (k == # bits in exp)
 - $\text{Bias} = 2^{k-1} - 1$
 - Single precision: 127 (so exp : 1...254, E : -126...127)
 - Double precision: 1023 (so exp : 1...2046, E : -1022...1023)
 - These enable negative values for E , for representing very small values
- **Significand coded with implied leading 1:** $M = 1 . \text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: the n bits of frac
 - Minimum when 000...0 ($M = 1.0$)
 - Maximum when 111...1 ($M = 2.0 - \epsilon$)
 - Get extra leading bit for “free”

Normalized Encoding Example

$$V = (-1)^S * M * 2^E$$



■ Value: `float f = 12345.0;`

$$\begin{aligned} 12345_{10} &= 11000000111001_2 \\ &= 1.1000000111001_2 \times 2^{13} \quad (\text{normalized form}) \end{aligned}$$

■ Significand:

$$\begin{aligned} M &= 1.\underline{1000000111001}_2 \\ \text{frac} &= \underline{1000000111001}0000000000_2 \end{aligned}$$

■ Exponent: $E = \text{exp} - \text{Bias}$, so $\text{exp} = E + \text{Bias}$

$$\begin{aligned} E &= 13 \\ \text{Bias} &= 127 \\ \text{exp} &= 140 = 10001100_2 \end{aligned}$$

■ Result:

