


Section 5: Arrays & Other Data Structures

- Array allocation and access in memory
- Multi-dimensional or nested arrays
- Multi-level arrays
- Other structures in memory
- Data structures and alignment

Structures & Alignment

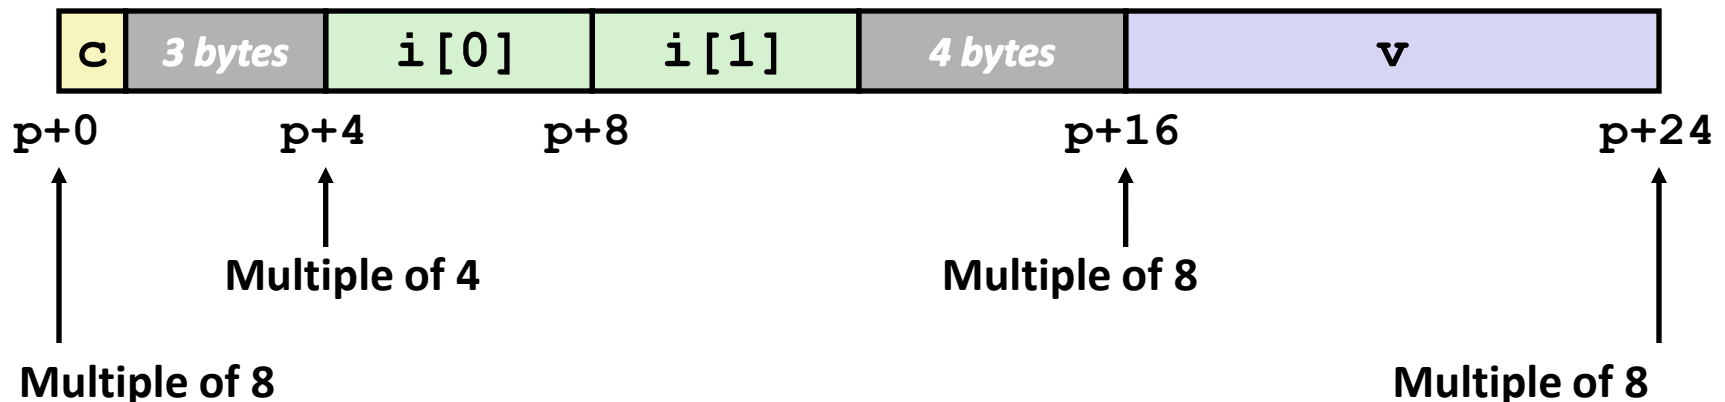
■ Unaligned Data



```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



Alignment Principles

- **Aligned Data**
 - Primitive data type requires K bytes
 - Address must be multiple of K
- **Aligned data is required on some machines; it is *advised* on IA32**
 - Treated differently by IA32 Linux, x86-64 Linux, and Windows!
- **What is the motivation for alignment?**

Alignment Principles

■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K

■ Aligned data is required on some machines; it is *advised* on IA32

- Treated differently by IA32 Linux, x86-64 Linux, and Windows!

■ Motivation for Aligning Data

- Physical memory is accessed by aligned chunks of 4 or 8 bytes (system-dependent)
 - Inefficient to load or store datum that spans quad word boundaries
- Also, virtual memory is very tricky when datum spans two pages (later...)

■ Compiler

- Inserts padding in structure to ensure correct alignment of fields
- `sizeof()` should be used to get true size of structs

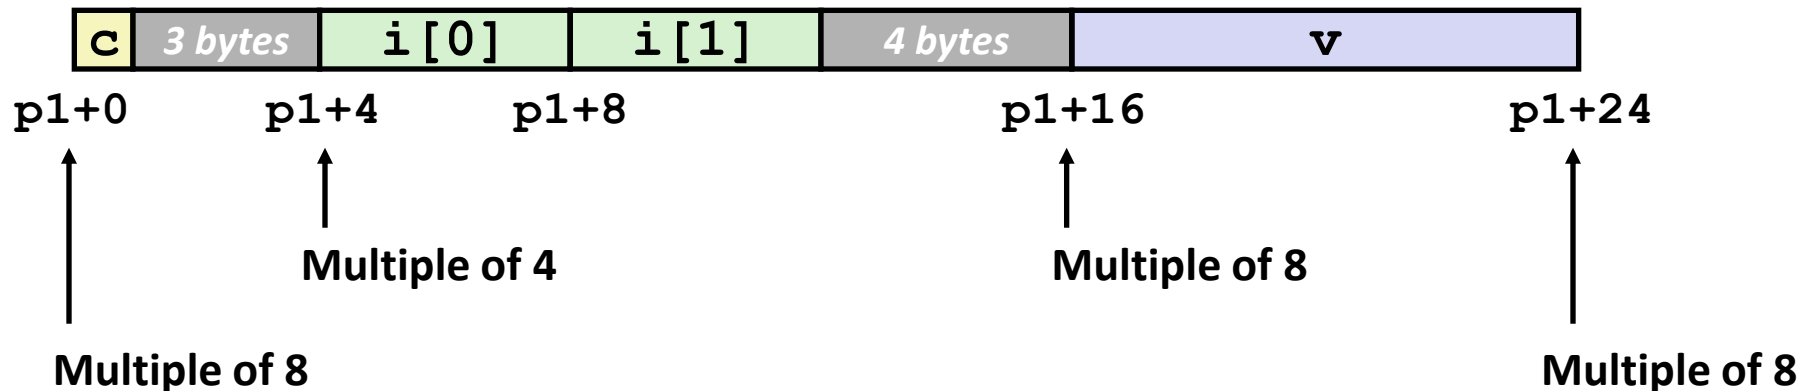
Specific Cases of Alignment (IA32)

- **1 byte: char, ...**
 - no restrictions on address
- **2 bytes: short, ...**
 - lowest 1 bit of address must be 0_2
- **4 bytes: int, float, char *, ...**
 - lowest 2 bits of address must be 00_2
- **8 bytes: double, ...**
 - Windows (and most other OSs & instruction sets): lowest 3 bits 000_2
 - Linux: lowest 2 bits of address must be 00_2
 - i.e., treated the same as a 4-byte primitive data type
- **12 bytes: long double**
 - Windows, Linux: lowest 2 bits of address must be 00_2

Satisfying Alignment with Structures

- **Within structure:**
 - Must satisfy every member's alignment requirement
- **Overall structure placement**
 - Each structure has alignment requirement K
 - K = Largest alignment of any element
 - Initial address & structure length must be multiples of K
- **Example (under Windows or x86-64):** $K = ?$
 - $K = 8$, due to **double** member

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p1;
```




Different Alignment Conventions

■ IA32 Windows or x86-64:


- $K = 8$, due to **double** member

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p1;
```



■ IA32 Linux: $K = ?$

- $K = 4$; **double** aligned like a 4-byte data type



Saving Space



- Put large data types first:

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p1;
```



```
struct S2 {
    double v;
    int i[2];
    char c;
} *p2;
```

- Effect (example x86-64, both have $K=8$)

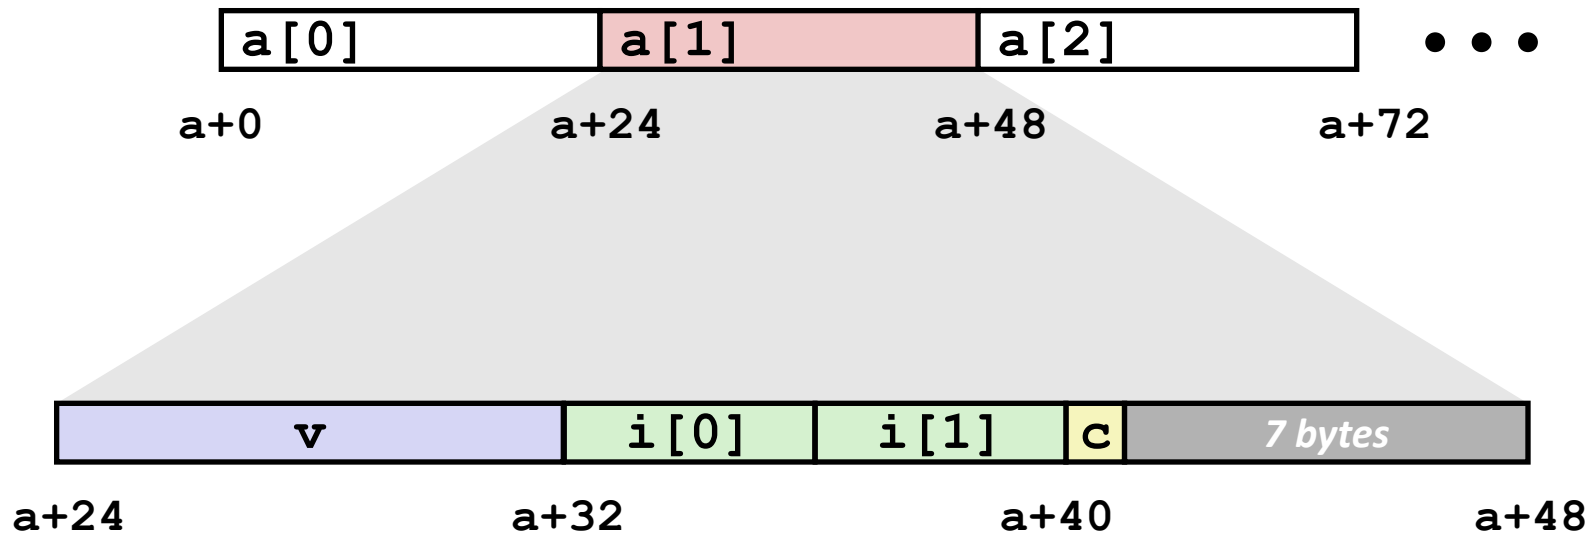


Unfortunately, doesn't satisfy requirement that struct's *total size* is a multiple of K

Arrays of Structures

- Satisfy alignment requirement for every element

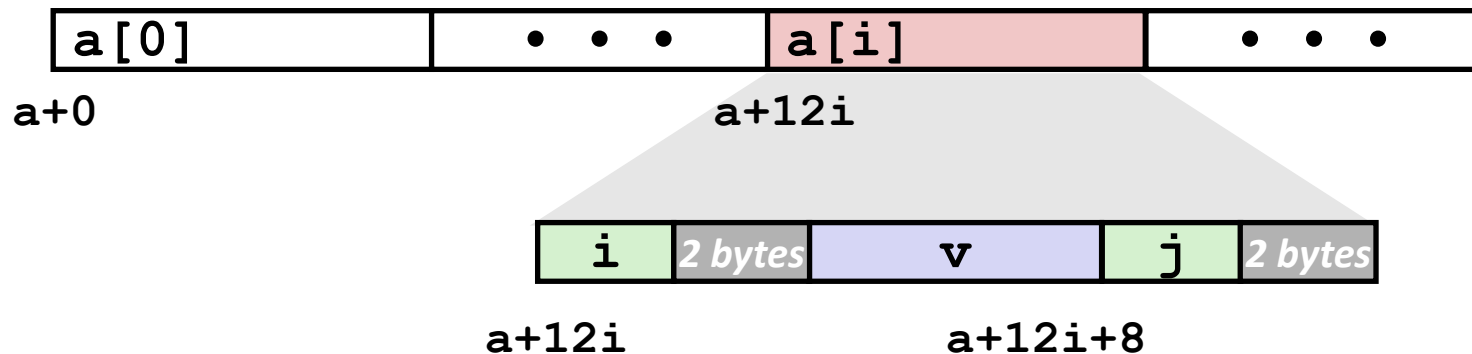
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Array Elements

- Compute array offset $12i$ (`sizeof(S3)`)
- Element j is at offset 8 within structure
- Since a is static array, assembler gives offset $a+8$

```
// Global:
struct S3 {
    short i;
    float v;
    short j;
} a[10];
```



```
short get_j(int idx)
{
    return a[idx].j;
// return (a + idx)->j;
}
```

```
# %eax = idx
leal (%eax,%eax,2),%eax # 3*idx
movswl a+8(,%eax,4),%eax # a+12*idx+8
```