

Section 8: Processes

- ~~What is a process~~
- ~~Creating processes~~
- Fork-Exec

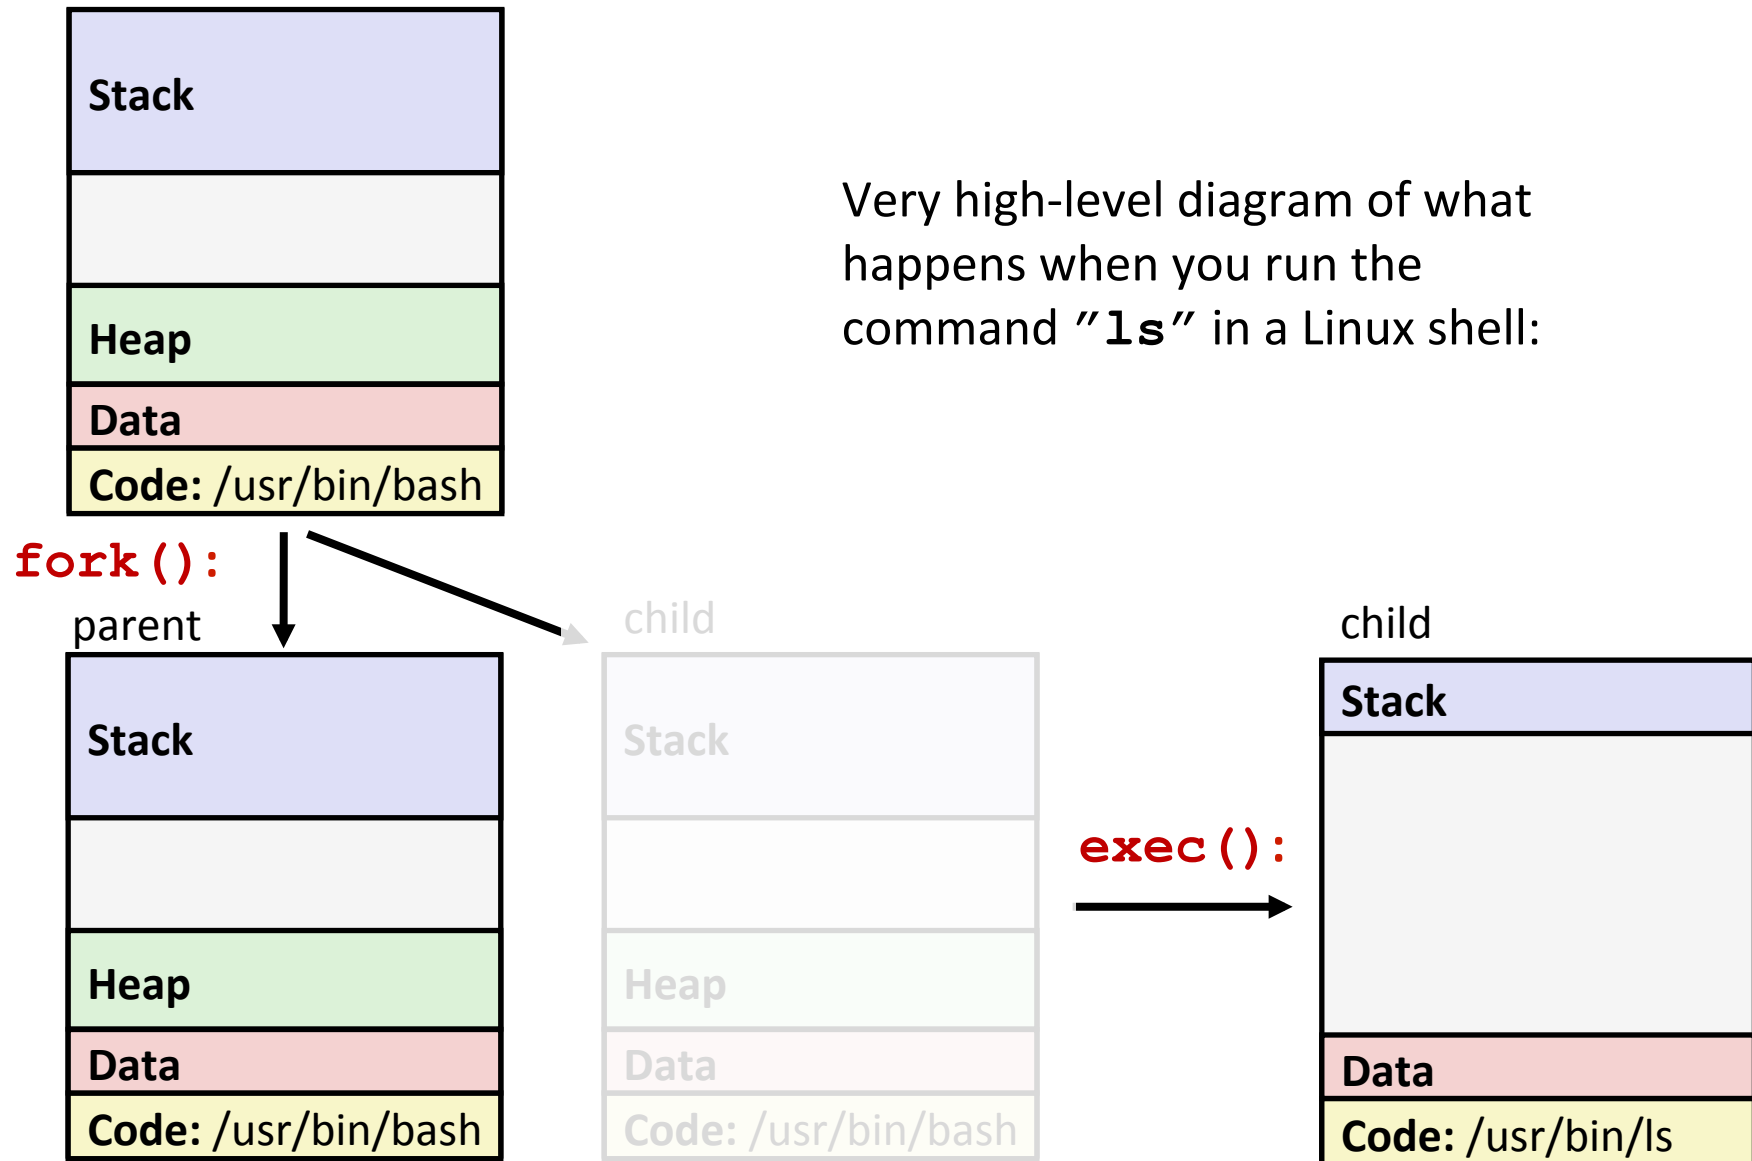
Fork-Exec

■ fork-exec model:

- **fork()** creates a copy of the current process
- **execve()** replaces the current process' code & address space with the code for a different program
 - There is a whole family of **exec** calls – see **exec(3)** and **execve(2)**

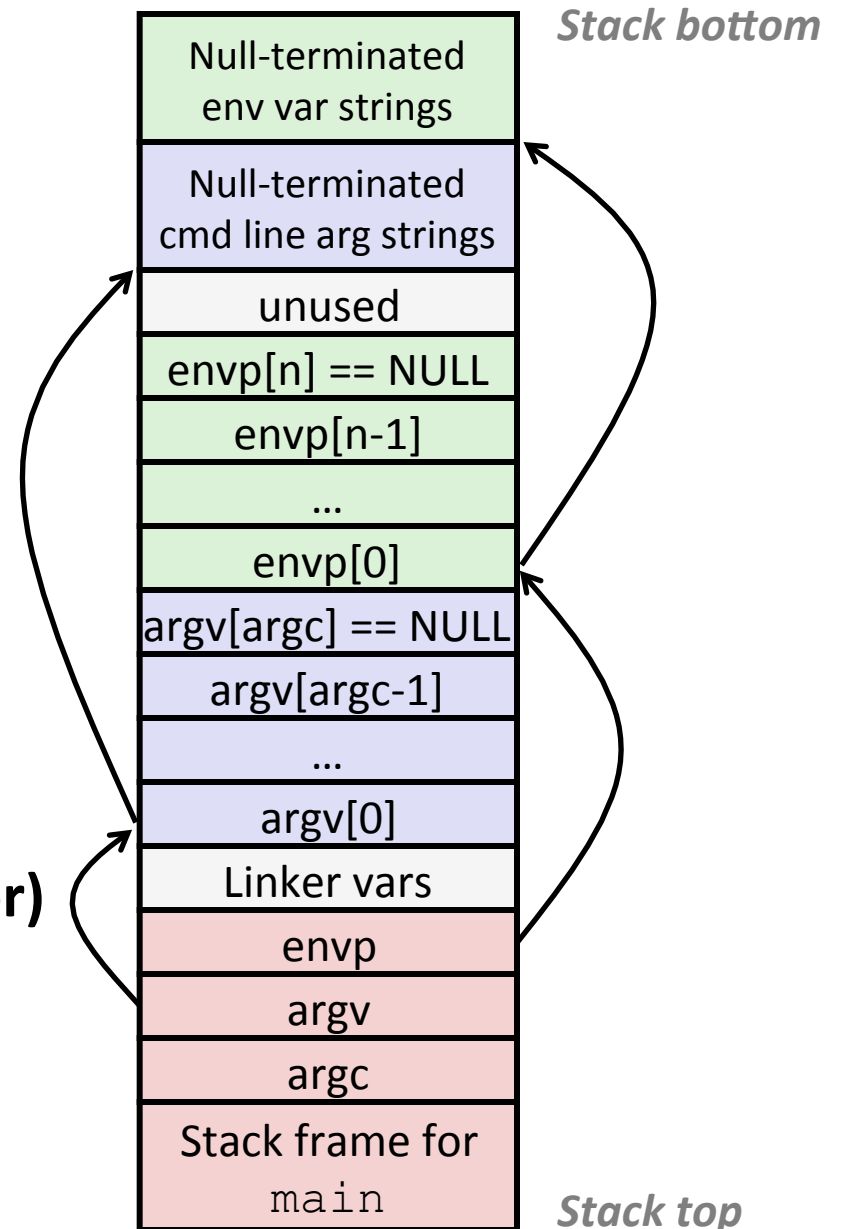
```
// Example arguments: path="/usr/bin/ls",  
//      argv[0]="/usr/bin/ls", argv[1]="-ahl", argv[2]=NULL  
void fork_exec(char *path, char *argv[])  
{  
    pid_t pid = fork();  
    if (pid != 0) {  
        printf("Parent: created a child %d\n", pid);  
    } else {  
        printf("Child: exec-ing new program now\n");  
        execv(path, argv);  
    }  
    printf("This line printed by parent only!\n");  
}
```

Exec-ing a new program



execve: Loading and Running Programs

- `int execve(`
`char *filename,`
`char *argv[],`
`char *envp[]`
`)`
- **Loads and runs in current process:**
 - Executable `filename`
 - With argument list `argv`
 - And environment variable list `envp`
 - Env. vars: “name=value” strings
(e.g. “PWD=/homes/iws/pjh”)
- ***execve does not return* (unless error)**
- **Overwrites code, data, and stack**
 - Keeps pid, open files, a few other items



exit: Ending a process

■ `void exit(int status)`

- Exits a process
 - Status code: 0 is used for a normal exit, nonzero for abnormal exit
- `atexit()` registers functions to be executed upon exit

```
void cleanup(void) {  
    printf("cleaning up\n");  
}  
  
void fork6() {  
    atexit(cleanup);  
    fork();  
    exit(0);  
}
```

Zombies



■ Idea

- When process terminates, it still consumes system resources
 - Various tables maintained by OS
- Called a “zombie”
 - A living corpse, half alive and half dead

■ Reaping

- Performed by parent on terminated child
- Parent is given exit status information
- Kernel discards process

■ What if parent doesn't reap?


- If any parent terminates without reaping a child, then child will be reaped by **init** process (`pid == 1`)
- But in long-running processes we need *explicit* reaping
 - e.g., shells and servers

`wait`: Synchronizing with Children

- `int wait(int *child_status)`
 - Suspends current process (i.e. the parent) until one of its children terminates
 - Return value is the `pid` of the child process that terminated
 - On successful return, the child process is reaped
 - If `child_status != NULL`, then the `int` that it points to will be set to a status indicating why the child process terminated
 - There are special macros for interpreting this status – see `wait(2)`
- If parent process has multiple children, `wait()` will return when *any* of the children terminates
 - `waitpid()` can be used to wait on a specific child process

wait Example

```
void fork_wait() {  
    int child_status;  
    pid_t child_pid;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
    } else {  
        child_pid = wait(&child_status);  
        printf("CT: child %d has terminated\n",  
              child_pid);  
    }  
    printf("Bye\n");  
    exit(0);  
}
```



Process management summary

- **`fork` gets us two copies of the same process (but `fork()` returns different values to the two processes)**
- **`execve` has a new process substitute itself for the one that called it**
 - Two-process program:
 - First **`fork()`**
 - `if (pid == 0) { //child code } else { //parent code }`
 - Two different programs:
 - First **`fork()`**
 - `if (pid == 0) { execve() } else { //parent code }`
 - Now running two completely different programs
- **`wait / waitpid` used to synchronize parent/child execution and to reap child process**

Summary

■ Processes

- At any given time, system has multiple active processes
- Only one can execute at a time, but each process appears to have total control of the processor
- OS periodically “context switches” between active processes
 - Implemented using *exceptional control flow*

■ Process management

- fork-exec model