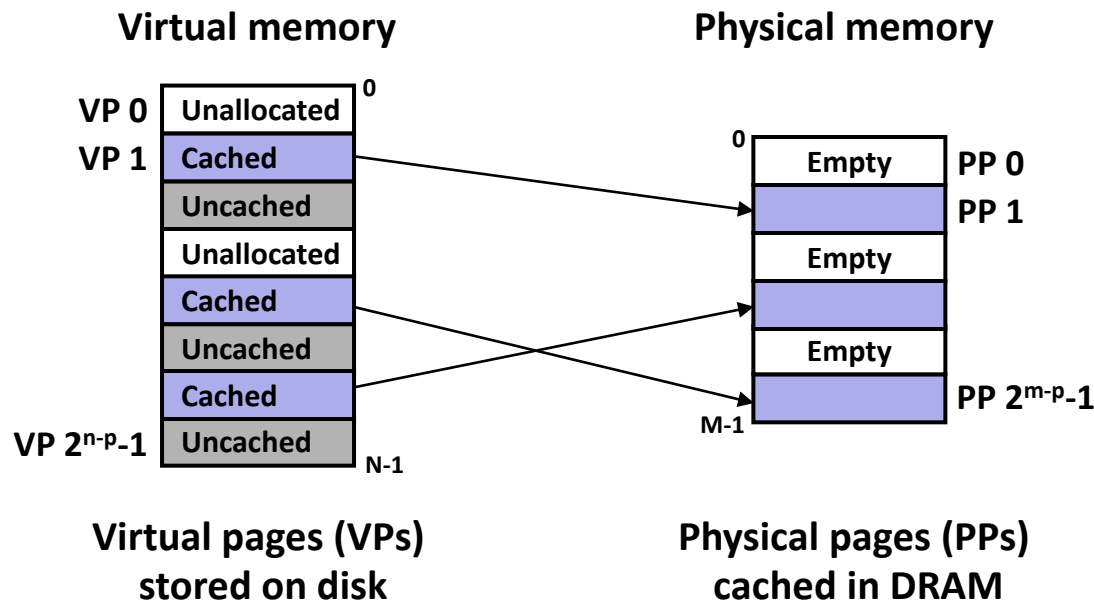


# Section 9: Virtual Memory (VM)

- Overview and motivation
- Indirection
- VM as a tool for caching
- Memory management/protection and address translation
- Virtual memory example

# VM and the Memory Hierarchy

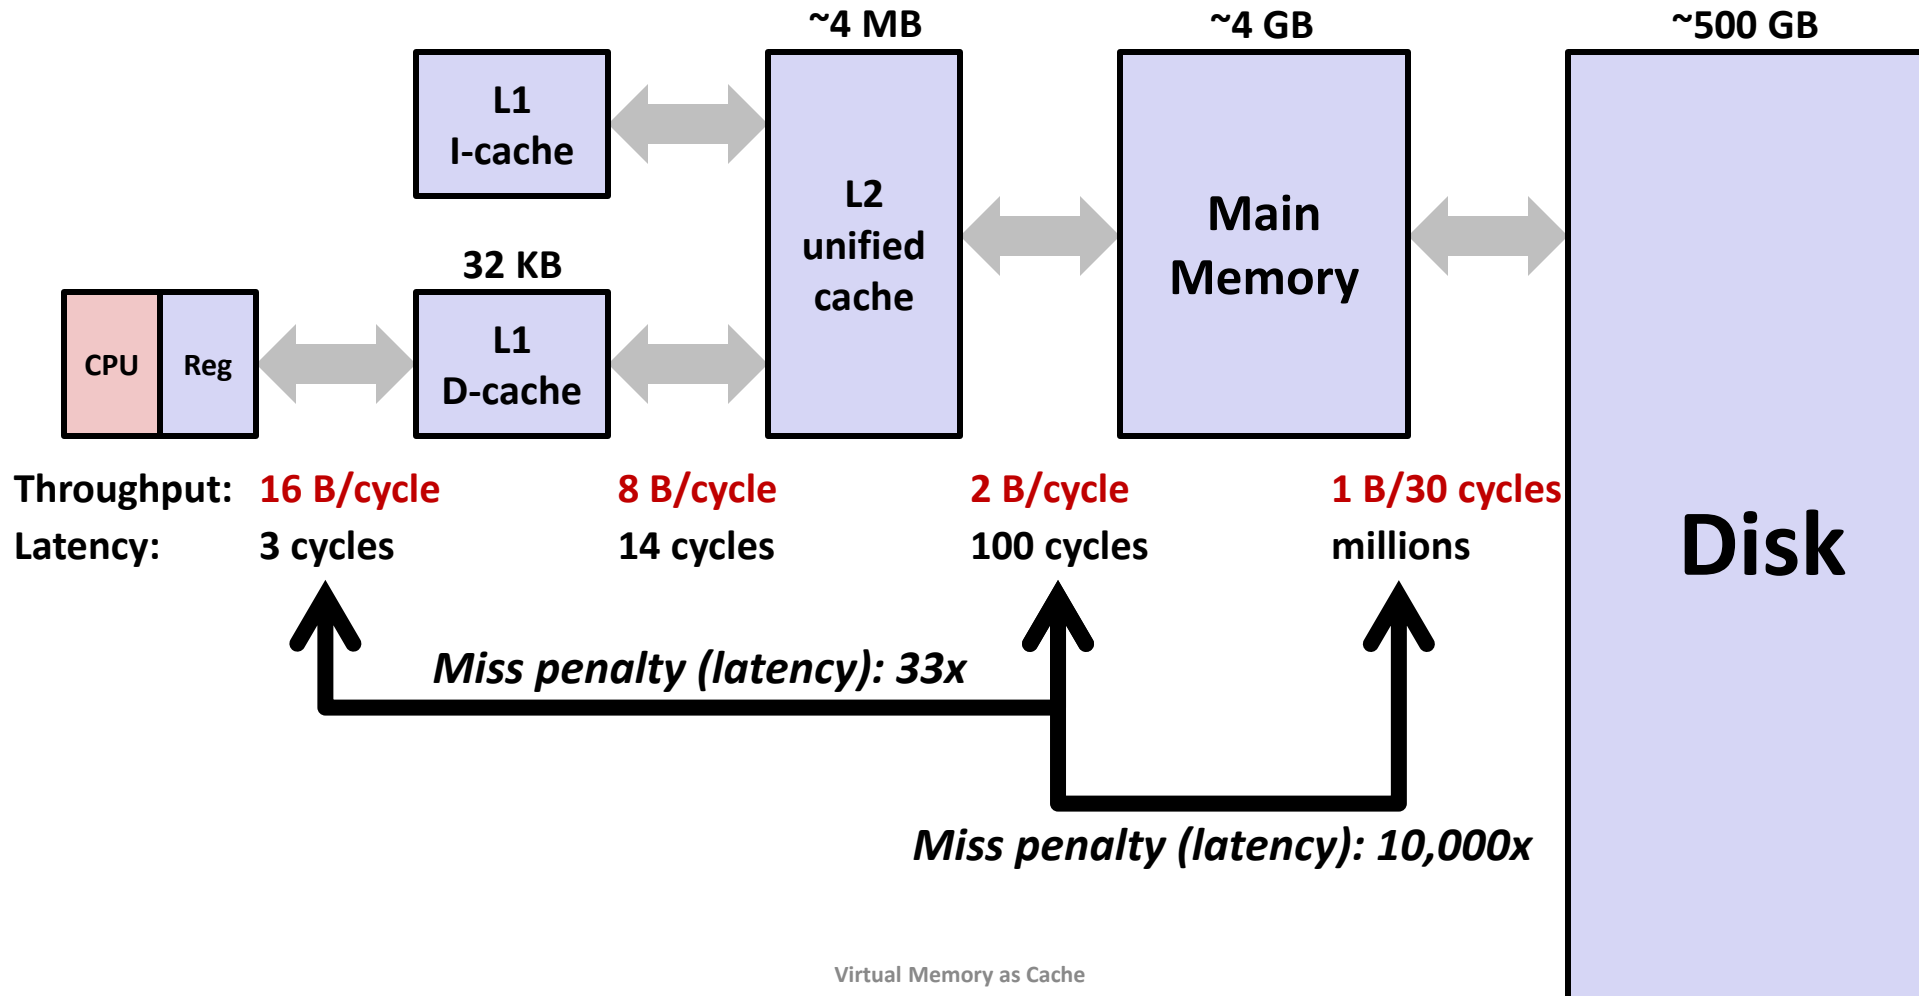
- Think of virtual memory as an array of  $N = 2^n$  contiguous bytes stored *on a disk*
- Then physical main memory (DRAM) is used as a *cache* for the virtual memory array
  - The cache blocks are called *pages* (size is  $P = 2^p$  bytes)



# Memory Hierarchy: Core 2 Duo

*Not drawn to scale*

L1/L2 cache: 64 B blocks



# DRAM Cache Organization

- **DRAM cache organization driven by the enormous miss penalty**
  - DRAM is about **10x** slower than SRAM
  - Disk is about **10,000x** slower than DRAM
    - (for first byte; faster for next byte)
- **Consequences?**
  - Block size?
  - Associativity?
  - Write-through or write-back?

# DRAM Cache Organization

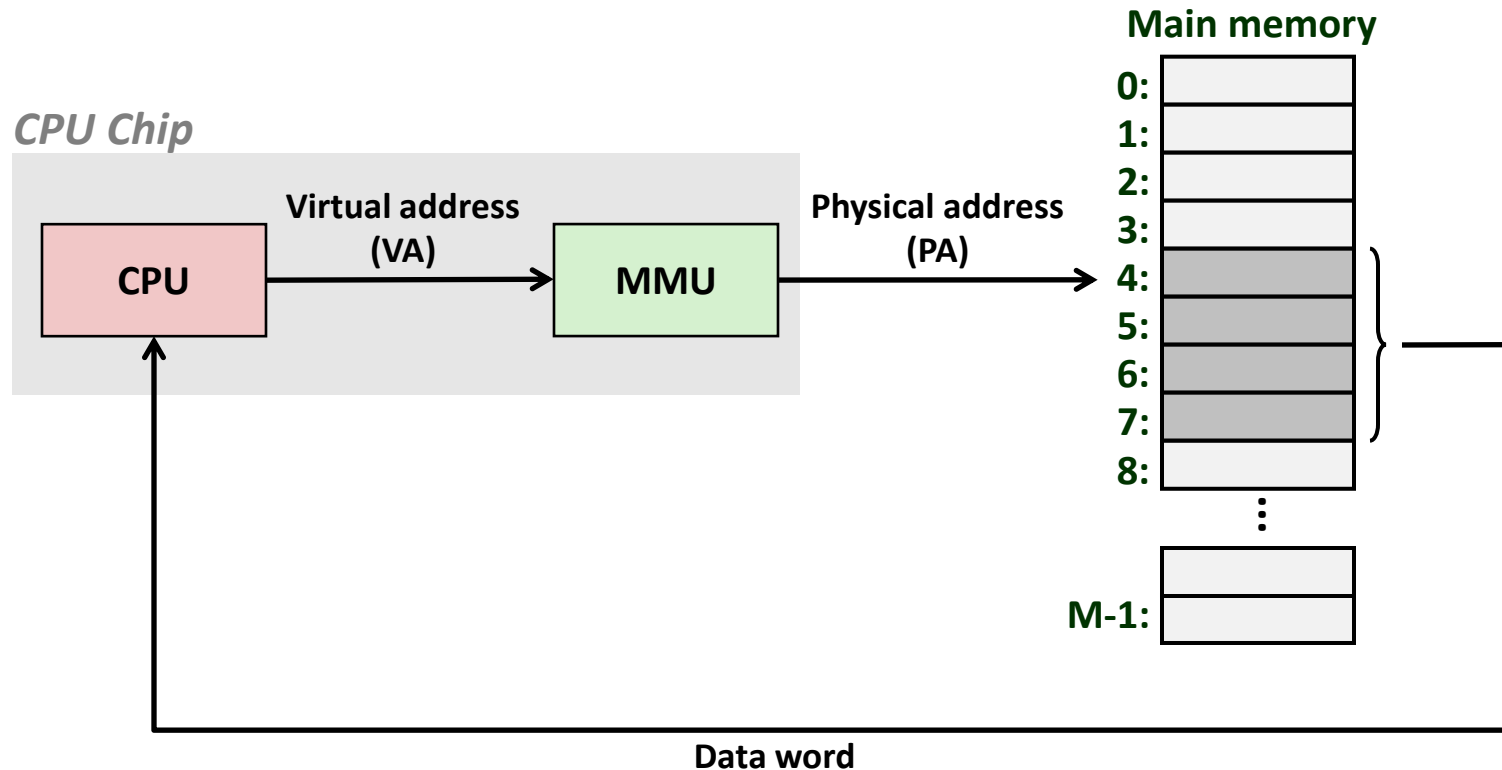
## ■ DRAM cache organization driven by the enormous miss penalty

- DRAM is about 10x slower than SRAM
- Disk is about 10,000x slower than DRAM
  - (for first byte; faster for next byte)

## ■ Consequences

- Large page (block) size: typically 4-8 KB, sometimes 4 MB
- Fully associative
  - Any VP can be placed in any PP
  - Requires a “large” mapping function – different from CPU caches
- Highly sophisticated, expensive replacement algorithms
  - Too complicated and open-ended to be implemented in hardware
- Write-back rather than write-through

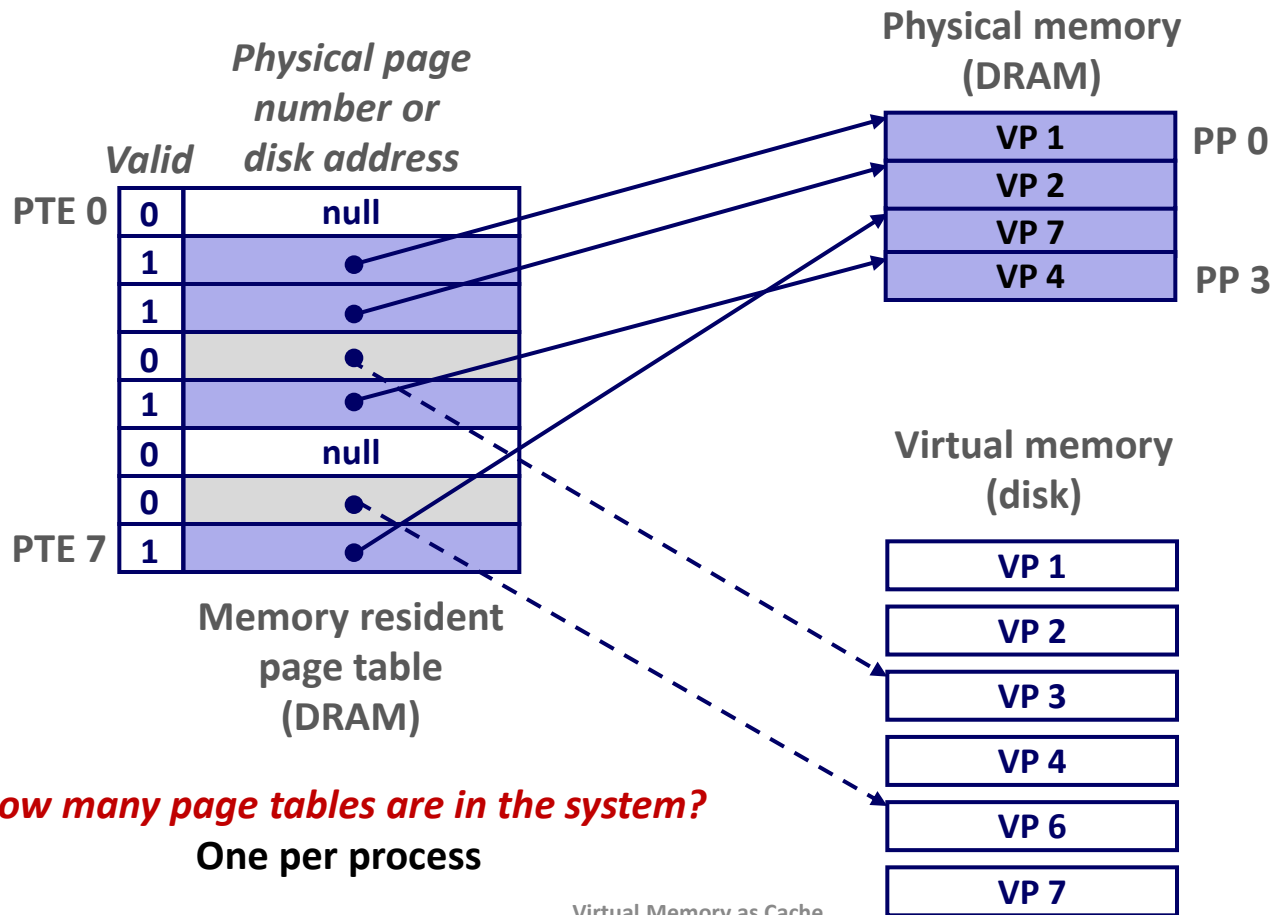
# Indexing into the “DRAM Cache”



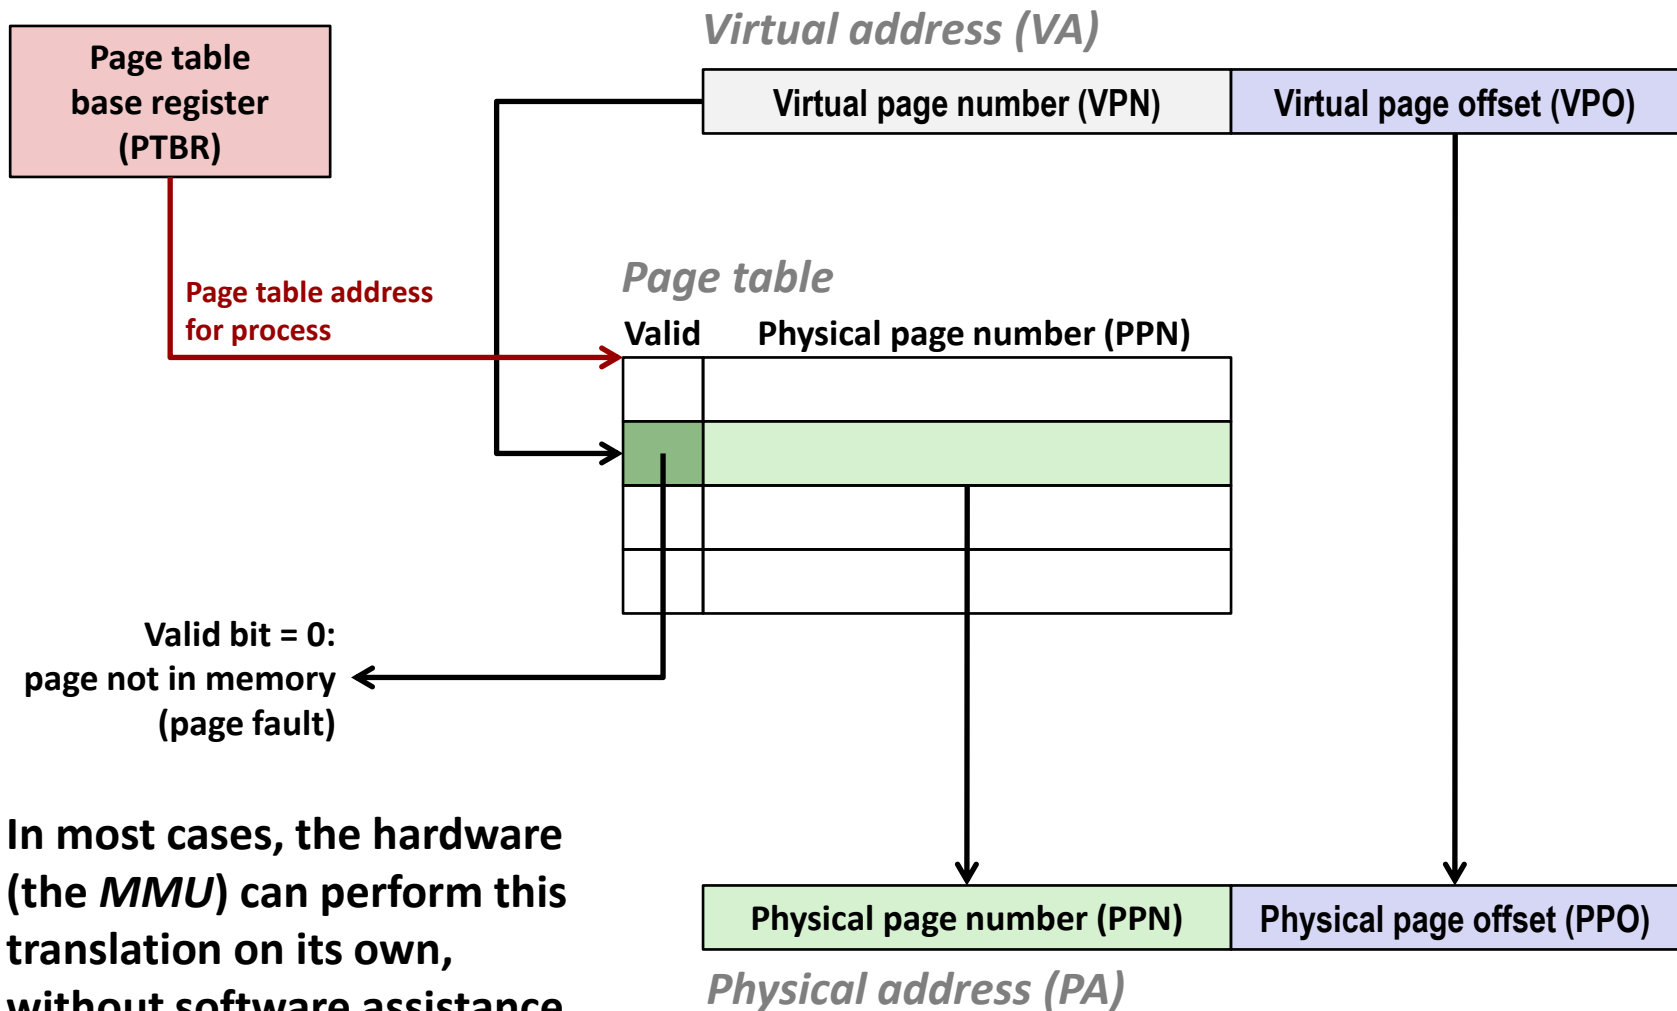
*How do we perform the VA -> PA translation?*

# Address Translation: Page Tables

- A **page table** (PT) is an array of **page table entries** (PTEs) that maps virtual pages to physical pages.



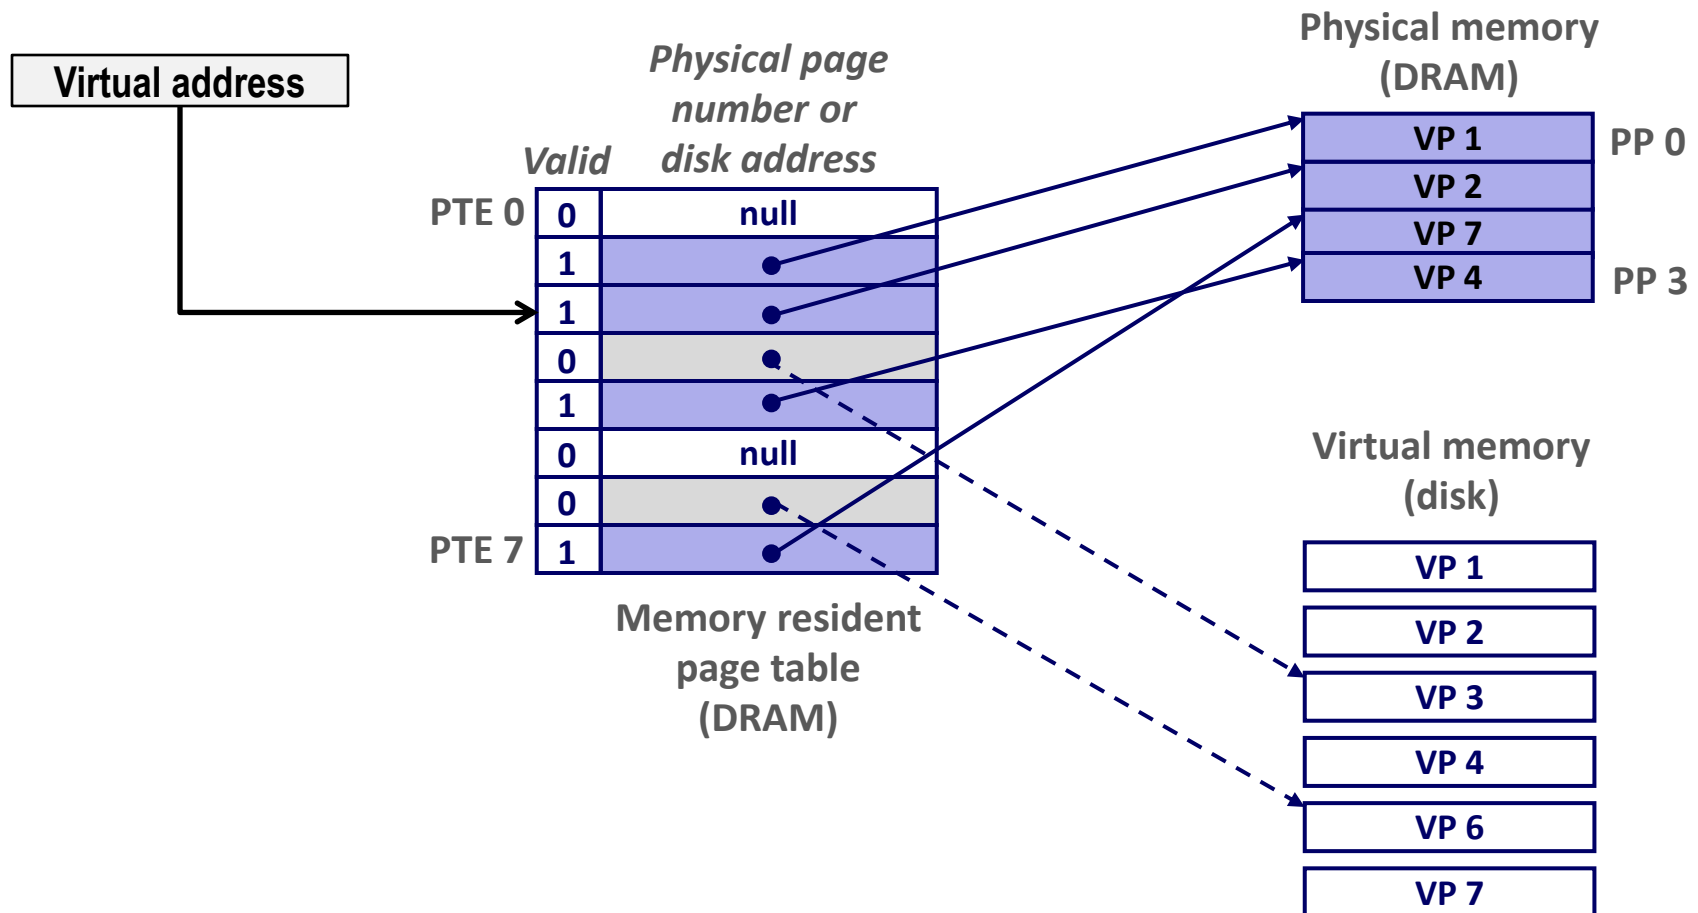
# Address Translation With a Page Table





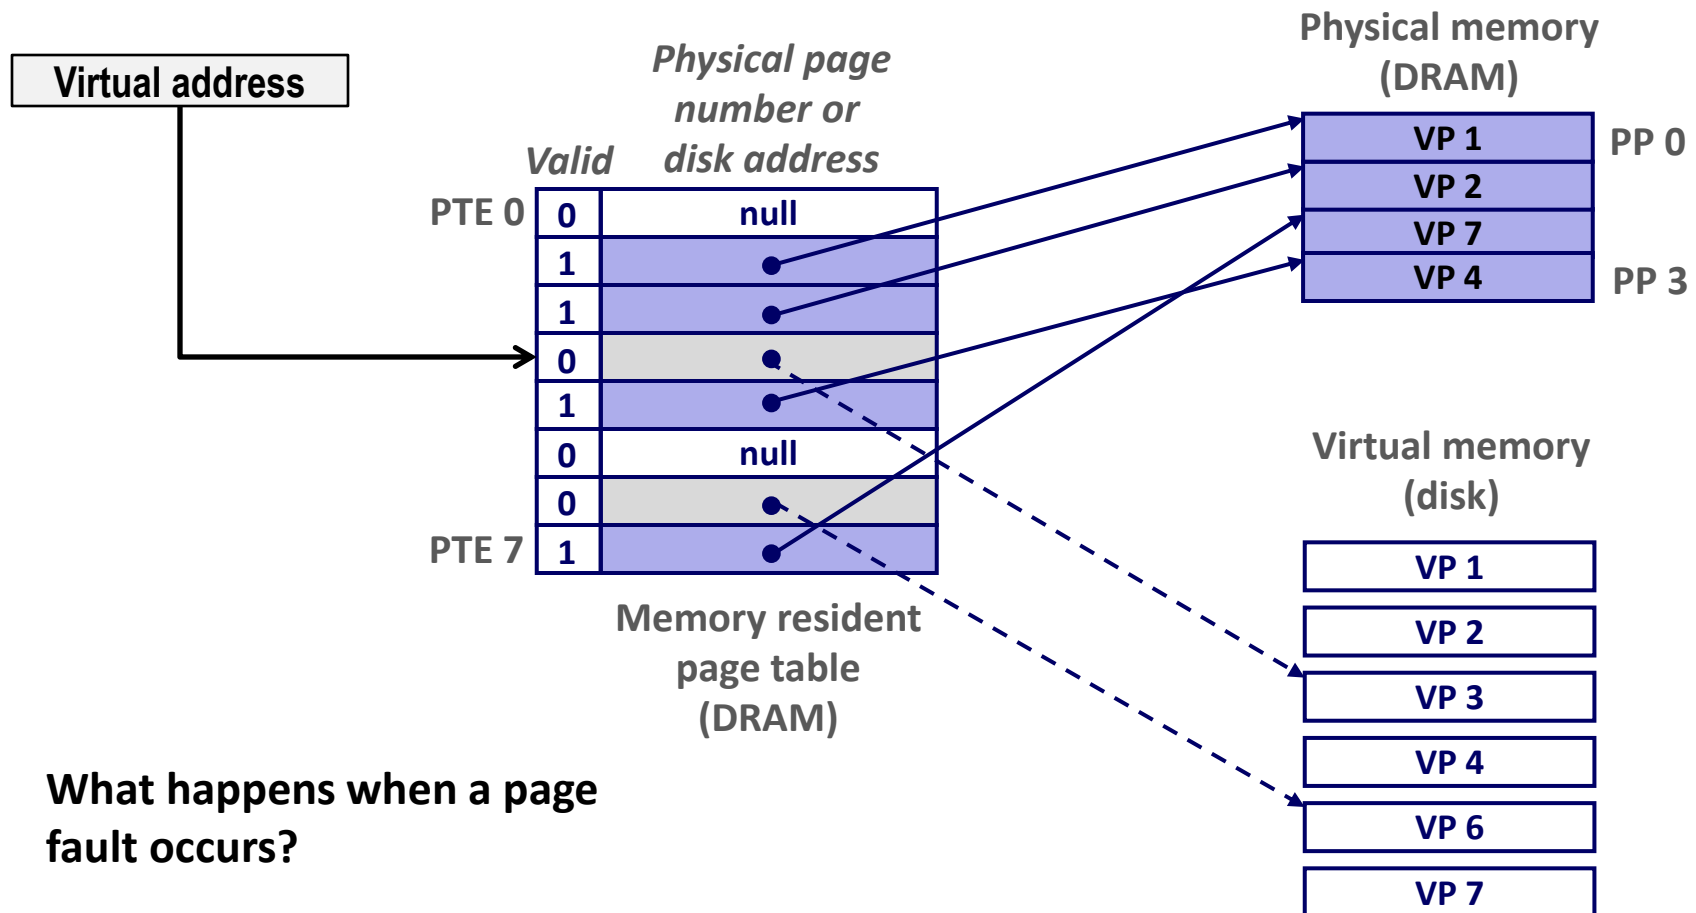
# Page Hit

- **Page hit:** reference to VM byte that is in physical memory



# Page Fault

- **Page fault:** reference to VM byte that is **NOT** in physical memory



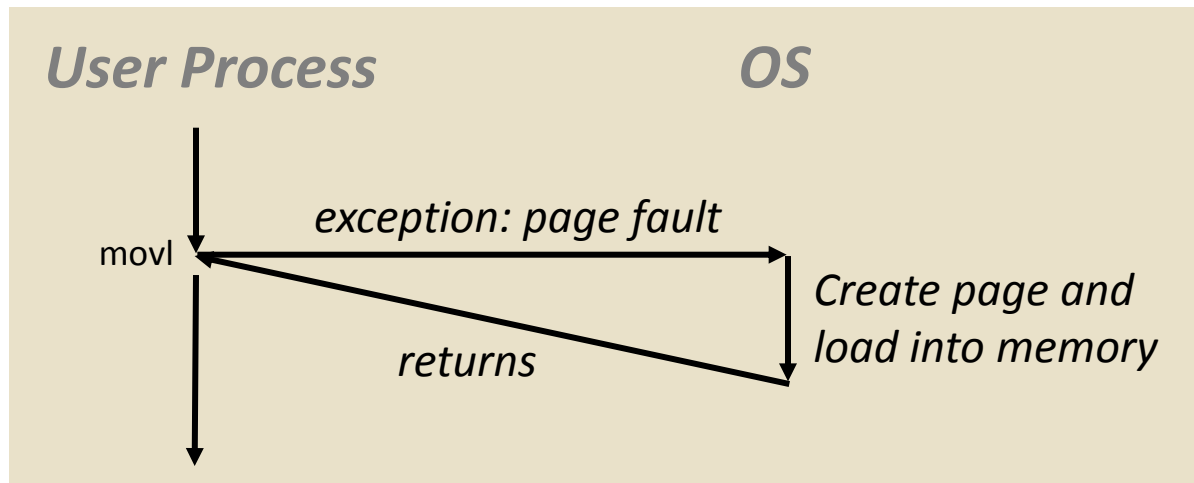
What happens when a page fault occurs?

# Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

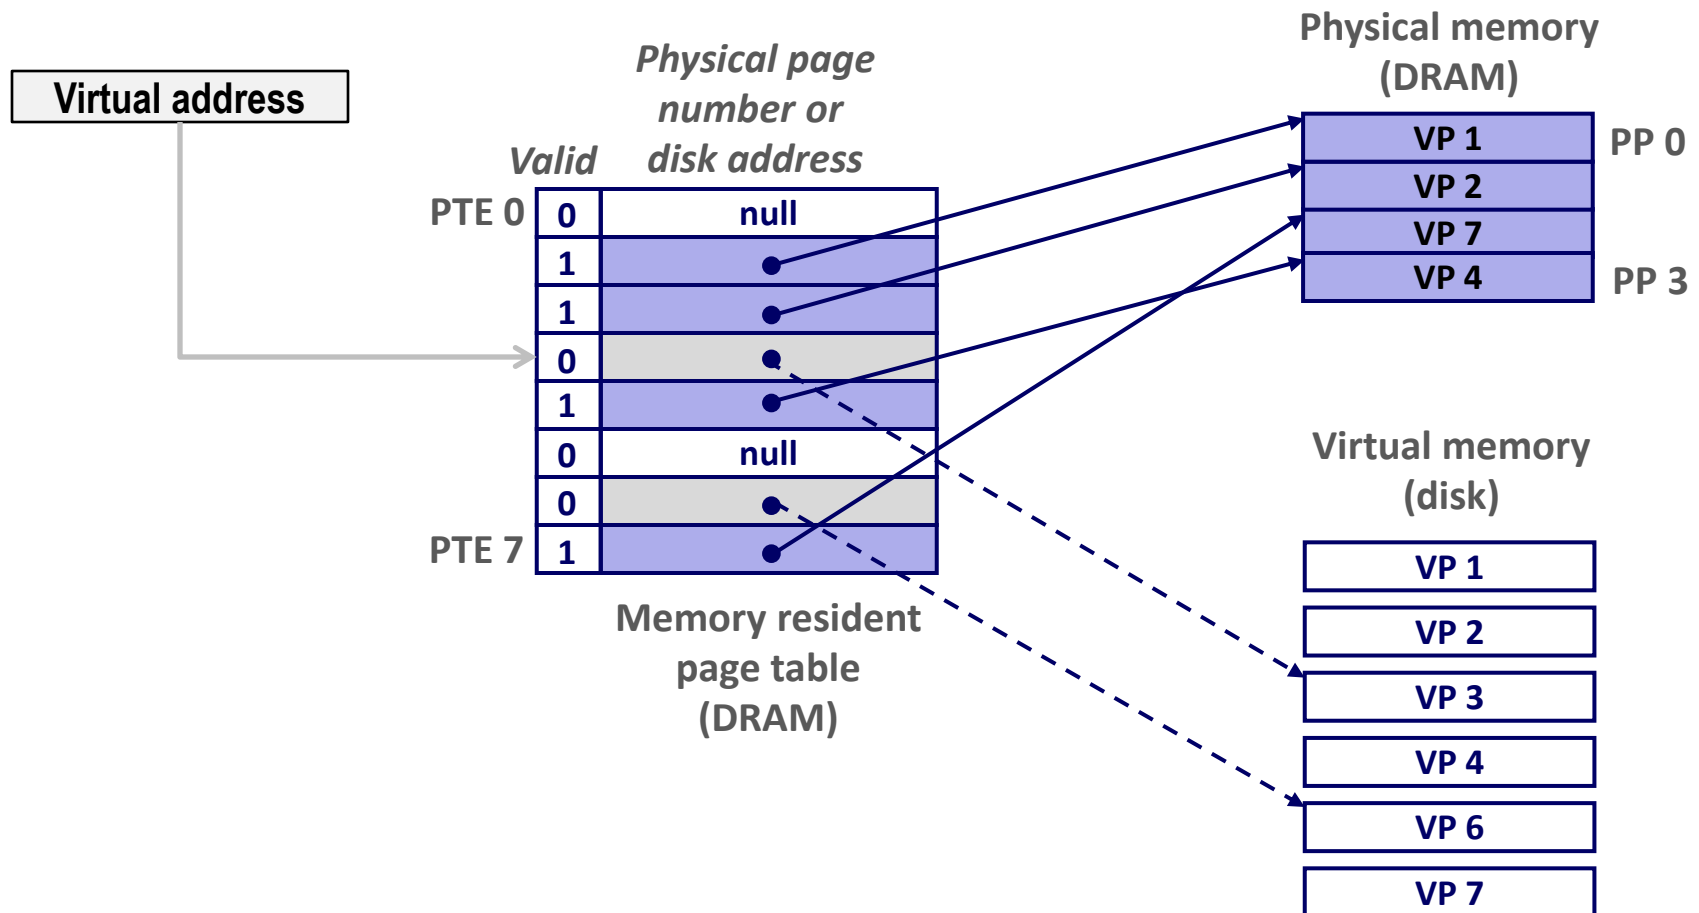
80483b7:	c7 05 10 9d 04 08 0d	movl	\$0xd,0x8049d10
----------	----------------------	------	-----------------



- Page handler must load page into physical memory
- Returns to faulting instruction: **mov** is executed again!
- Successful on second try

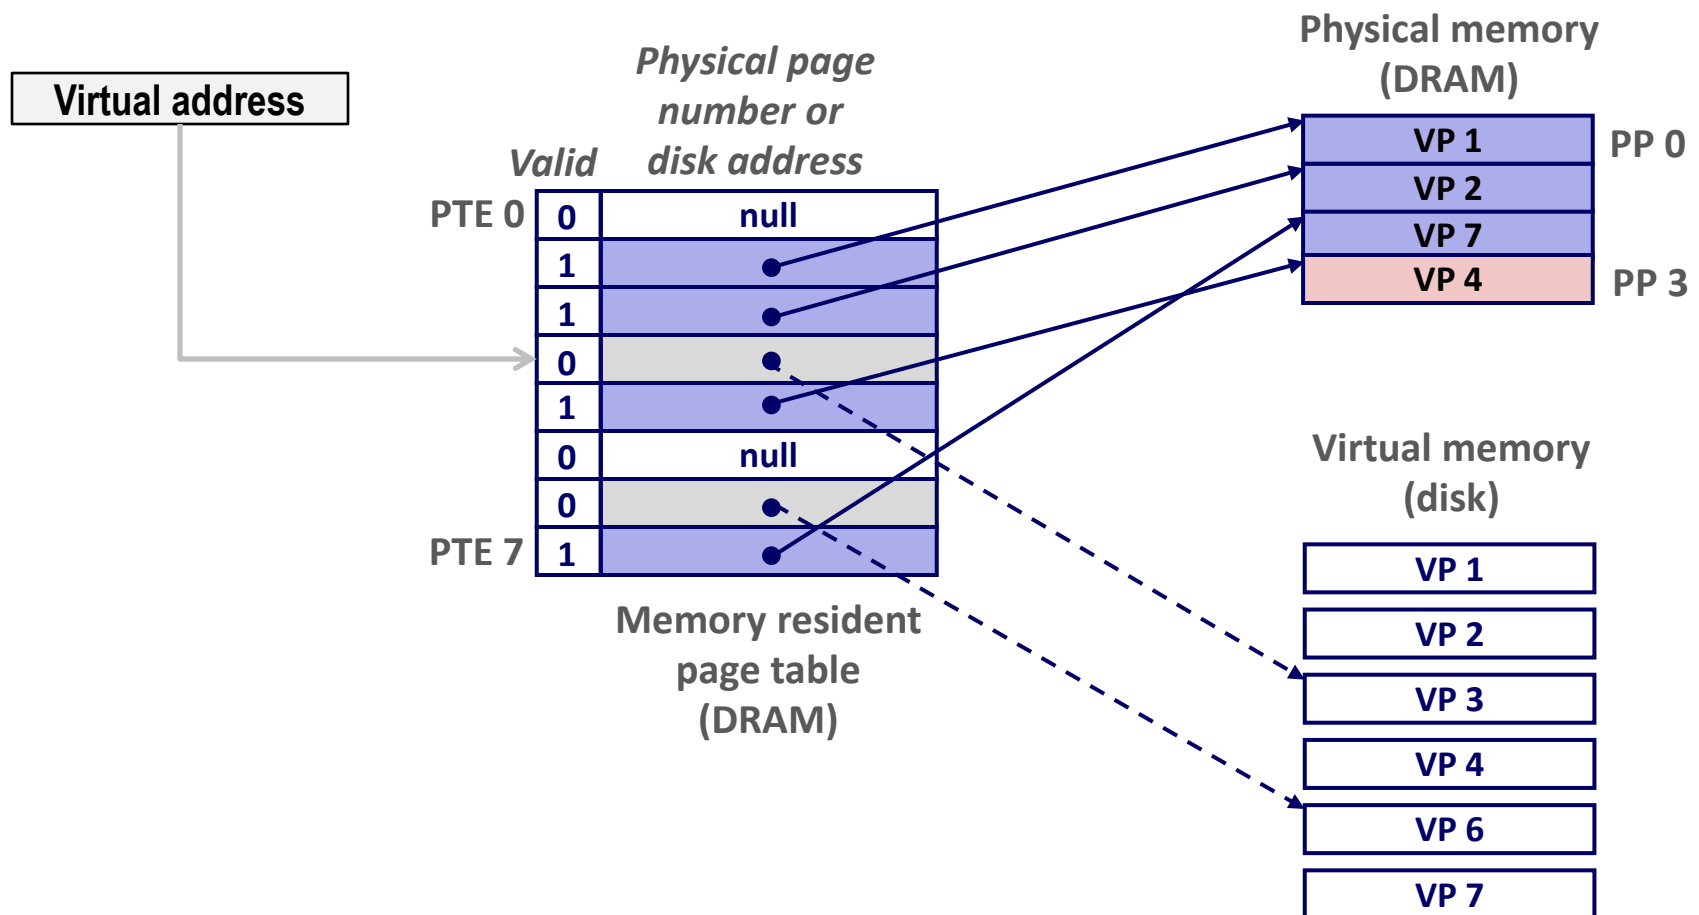
# Handling Page Fault

- Page miss causes page fault (an exception)



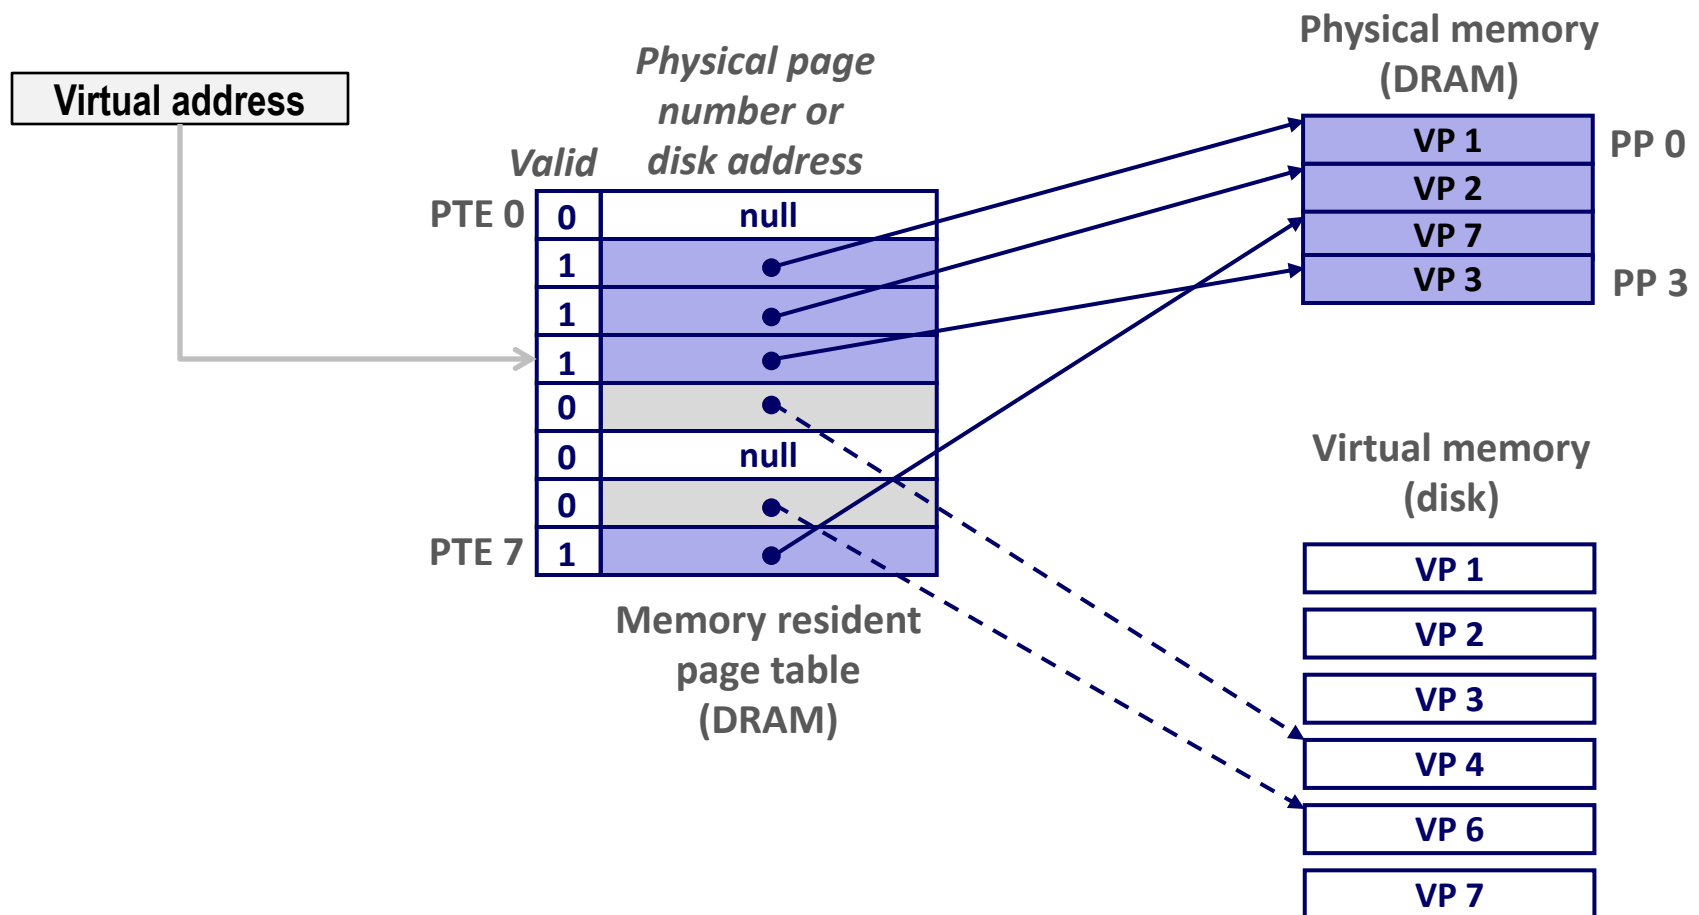
# Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a *victim* to be evicted (here VP 4)



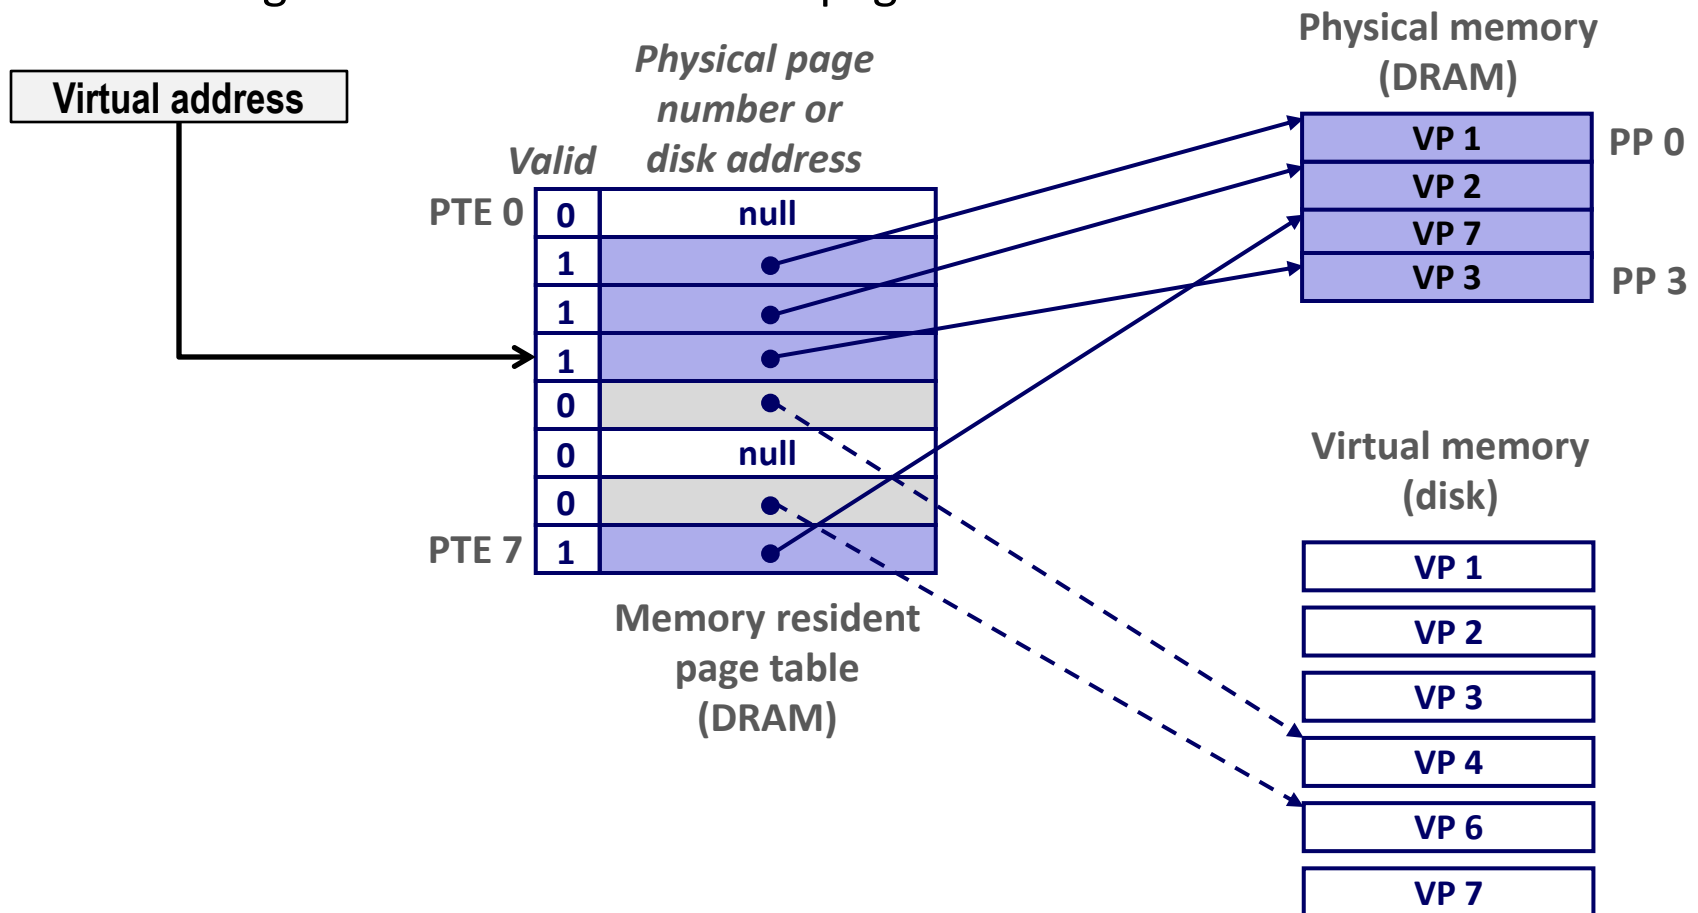
# Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a *victim* to be evicted (here VP 4)



# Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a *victim* to be evicted (here VP 4)
- Offending instruction is restarted: page hit!



# Why does it work? Locality

- **Virtual memory works well because of locality**
  - Same reason that L1 / L2 / L3 caches work
- **The set of virtual pages that a program is “actively” accessing at any point in time is called its *working set***
  - Programs with better temporal locality will have smaller working sets
- **If (working set size < main memory size):**
  - Good performance for one process after compulsory misses
- **If (SUM(working set sizes) > main memory size):**
  - *Thrashing*: Performance meltdown where pages are swapped (copied) in and out continuously