

Section 10: Memory Allocation Topics

■ Dynamic memory allocation

- Size/number of data structures may only be known at run time
- Need to allocate space on the heap
- Need to de-allocate (free) unused memory so it can be re-allocated

■ Implementation


- Implicit free lists
- Explicit free lists – subject of next programming assignment
- Segregated free lists

■ Garbage collection

■ Common memory-related bugs in C programs


Assumptions Made in This Section

- **Memory is word addressed (each word can hold a pointer)**
 - block size is a multiple of words




Allocation Example


```
p1 = malloc(4)
```



```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(2)
```



What information does the allocator need to keep track of?

Constraints

■ Applications

- Can issue arbitrary sequence of malloc() and free() requests
- free() requests must be made only for a previously malloc()'d block

■ Allocators

- Can't control number or size of allocated blocks
- Must respond immediately to malloc() requests
 - *i.e.*, can't reorder or buffer requests
- Must allocate blocks from free memory
 - *i.e.*, blocks can't overlap
- Must align blocks so they satisfy all alignment requirements
 - 8 byte alignment for GNU malloc (**libc** malloc) on Linux boxes
- Can't move the allocated blocks once they are malloc()'d
 - *i.e.*, compaction is not allowed. *Why not?*

Performance Goal: Throughput

- Given some sequence of **malloc** and **free** requests:
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- Goals: maximize throughput and peak memory utilization
 - These goals are often conflicting
- Throughput:
 - Number of completed requests per unit time
 - Example:
 - 5,000 **malloc()** calls and 5,000 **free()** calls in 10 seconds
 - Throughput is 1,000 operations/second

Performance Goal: Peak Memory Utilization


- Given some sequence of `malloc` and `free` requests:
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- **Def: Aggregate payload P_k**
 - `malloc(p)` results in a block with a **payload** of p bytes
 - After request R_k has completed, the **aggregate payload** P_k is the sum of currently allocated payloads
- **Def: Current heap size = H_k**
 - Assume H_k is monotonically nondecreasing
 - Allocator can increase size of heap using `sbrk()`
- **Def: Peak memory utilization after k requests**
 - $U_k = (\max_{i \leq k} P_i) / H_k$
 - Goal: maximize utilization for a sequence of requests.
 - *Why is this hard? And what happens to throughput?*

Fragmentation

- Poor memory utilization is caused by *fragmentation*
 - *internal* fragmentation
 - *external* fragmentation

Internal Fragmentation

- For a given block, *internal fragmentation* occurs if payload is smaller than block size



- **Caused by**
 - overhead of maintaining heap data structures (inside block, outside payload)
 - padding for alignment purposes
 - explicit policy decisions (e.g., to return a big block to satisfy a small request)
why would anyone do that?
- **Depends only on the pattern of *previous* requests**
 - thus, easy to measure


External Fragmentation

- Occurs when there is enough aggregate heap memory, but no single free block is large enough


```
p1 = malloc(4)
```




```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(6)
```

Oops! (what would happen now?)

- Depends on the pattern of future requests
 - Thus, difficult to measure