

Section 10: Memory Allocation Topics

■ ~~Dynamic memory allocation~~

- ~~Size/number of data structures may only be known at run time~~
- ~~Need to allocate space on the heap~~
- ~~Need to de-allocate (free) unused memory so it can be re-allocated~~

■ Implementation

- Implicit free lists
- Explicit free lists – subject of next programming assignment
- Segregated free lists

■ Garbage collection

■ Common memory-related bugs in C programs


Implementation Issues

- How do we know how much memory to free given just a pointer?
- How do we keep track of the free blocks?
- How do we pick a block to use for allocation (when many might fit)?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
- How do we reinsert freed block into the heap?

Knowing How Much to Free


■ Standard method

- Keep the length of a block in the word preceding the block
 - This word is often called the *header field* or *header*
- Requires an extra word for every allocated block




Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks



- Method 2: *Explicit list* among the free blocks using pointers



- Method 3: *Segregated free list*
 - Different free lists for different size classes
- Method 4: *Blocks sorted by size*
 - Can use a balanced binary tree (e.g. red-black tree) with pointers within each free block, and the length used as a key

Implicit Free Lists

■ For each block we need: size, is-allocated?

- Could store this information in two words: wasteful!

■ Standard trick

- If blocks are aligned, some low-order size bits are always 0
- Instead of storing an always-0 bit, use it as a allocated/free flag
- When reading size, must remember to mask out this bit

e.g. with 8-byte alignment,
sizes look like:

00000000


00001000

00010000

00011000

...

*Format of
allocated and
free blocks*



a = 1: allocated block


a = 0: free block

size: block size

**payload: application data
(allocated blocks only)**

Implicit Free List Example

Sequence of blocks in heap: 2/0, 4/1, 8/0, 4/1 (size/allocated)



- **8-byte alignment**
 - May require initial unused word
 - Causes some internal fragmentation
- **One word (0/1) to mark end of list**

Implicit List: Finding a Free Block

*p gets the block *header*
 *p & 1 extracts the allocated bit
 *p & -2 masks the allocated bit, gets just the size

■ *First fit:*

- Search list from beginning, choose *first* free block that fits:

```
p = heap_start;
while ((p < end) &&          \\ not passed end
      ((*p & 1) ||          \\ already allocated
      (*p <= len)))         \\ too small
  p = p + (*p & -2);         \\ goto next block
```

- Can take time linear in total number of blocks (allocated and free)
- In practice it can cause “splinters” at beginning of list

■ *Next fit:*

- Like first-fit, but search list starting where previous search finished
- Should often be faster than first-fit: avoids re-scanning unhelpful blocks
- Some research suggests that fragmentation is worse


■ *Best fit:*

- Search the list, choose the *best* free block: fits, with fewest bytes left over
- Keeps fragments small—usually helps fragmentation
- Will typically run slower than first-fit


Implicit List: Allocating in Free Block

■ Allocating in a free block: *splitting*

- Since allocated space might be smaller than free space, we might want to split the block



`addblock(p, 4)`



```
void addblock(ptr p, int len) {
    int newsize = ((len + 1) >> 1) << 1; // round up to even
    int oldsize = *p & -2;                // mask out low bit
    *p = newsize | 1;                     // set new length + allocated
    if (newsize < oldsize)
        *(p+newsize) = oldsize - newsize; // set length in remaining
                                           // part of block
}
```



Implicit List: Freeing a Block

■ Simplest implementation:

- Need only clear the “allocated” flag

```
void free_block(ptr p) { *p = *p & -2 }
```

- But can lead to “false fragmentation”




`malloc(5)` ***Oops!***

There is enough free space, but the allocator won't be able to find it

Implicit List: Coalescing

- Join (*coalesce*) with next/previous blocks, if they are free
 - Coalescing with next block




```
void free_block(ptr p) {
    *p = *p & -2;           // clear allocated bit
    next = p + *p;          // find next block
    if ((*next & 1) == 0)
        *p = *p + *next;    // add to this block if
                             // not allocated
}
```

- But how do we coalesce with the *previous* block?

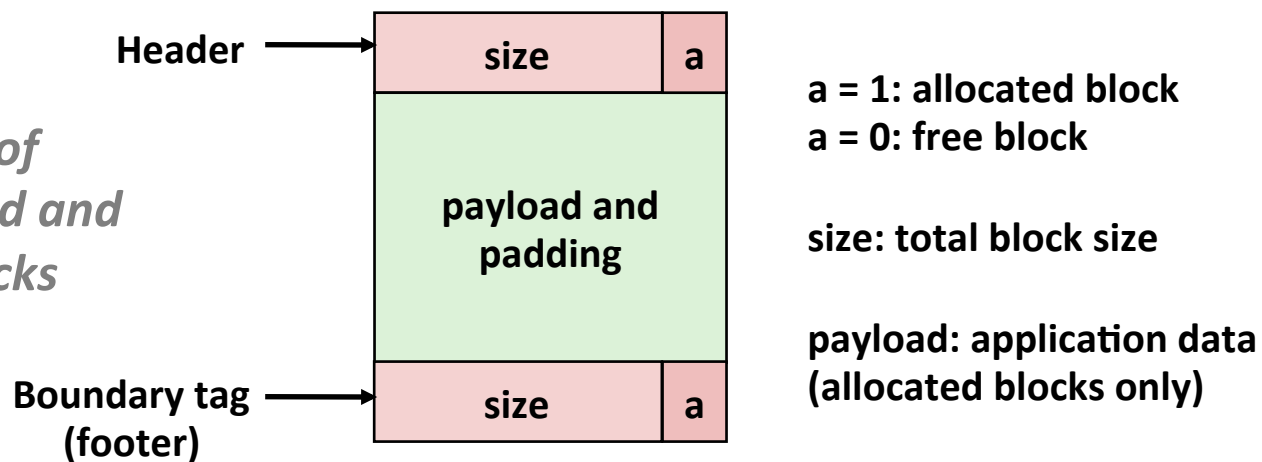
Implicit List: Bidirectional Coalescing

■ *Boundary tags* [Knuth73]

- Replicate size/allocated word at “bottom” (end) of free blocks
- Allows us to traverse the “list” backwards, but requires extra space
- Important and general technique!



*Format of
allocated and
free blocks*



Implicit Free Lists: Summary

- **Implementation: very simple**
- **Allocate cost:**
 - linear time (in total number of heap blocks) worst case
- **Free cost:**
 - constant time worst case
 - even with coalescing
- **Memory utilization:**
 - will depend on placement policy
 - First-fit, next-fit or best-fit
- **Not used in practice for `malloc()` / `free()` because of linear-time allocation**
 - used in some special purpose applications
- **The concepts of splitting and boundary tag coalescing are general to *all* allocators**